

This project guides you through the calculation of *PageRank*, the algorithm used by *Google* to rank search queries. The two graph files provided alongside this document (`tiny.txt`, `medium.txt`) represent two graphs of web-pages.

You will implement a model of search on the graph, the random crawler model, to compute the PageRank. You will do this in two ways: (i) By simulating the behaviour of a random crawler on the graph and (ii) By implementing a Markov Chain computation.

Each “page” on the web has links that can be followed to another page.

- 90% of the time, the crawler moves to another “page” by picking a **link** uniformly at random (with equal probability) from the current page and follows it.
- 10% of the time, the crawler picks a random page from the entire web and jumps to it. (Google originally used the values: 85% and 15% respectively, when PageRank was first introduced).

To complete the PageRank calculations, follow the steps laid out below:

- (1) You will need to read in one of the graph files provided (`tiny.txt`, `medium.txt`) as an *adjacency matrix*.

For this, create a function `adjMatrixFromFile(fileName)`. The function returns an *adjacency matrix* of the graph. The input graph files are given as a series of numbers, formatted in the following way:

- (a) The first line of the file contains a single integer, N . This value represents the number of nodes (i.e. web-pages) this graph contains. Store this value. It will be important when creating a properly sized matrix.
- (b) The file contains N more lines, one for each node in the graph. The lines have sets of integer pairs. Each integer pair (a , b) represents an edge in the graph, where there is an edge going from node a to node b . The node values will be in the range $[0, N-1]$. It is possible for the exact same pair to appear more than once.

An *adjacency matrix* is a square matrix containing values that show the distribution of edges within a graph. The adjacency matrix has dimensions $N \times N$. Each (row, column) entry: $a_{(i,j)}$ of the matrix contains the *number of edges* going $i \rightarrow j$. For example, location (1, 2) in the adjacency matrix generated from `tiny.txt` is 2, as the pair 1 2 occurs twice in the graph.

The `adjMatrixFromFile(fileName)` (the matrix read-in) function opens and reads a file `fileName`, creates an empty $N \times N$ matrix, then uses the contents of `fileName` to populate the adjacency matrix with data. You may find this easiest to do by reading the file line-by-line and using the `split()` function to break the line into a list of integers. Once the adjacency matrix is populated, close the file and return the adjacency matrix.

- (2) Next, create a function `outDegrees(adjacencyMatrix)` to generate an array of the *out-degree* of each node. The out-degree of a node is the number of edges which leave that node (*i.e.* the number of out-links in a web-page). This function accepts a single parameter, `adjacencyMatrix`, which is an $N \times N$ NumPy matrix. The out-degree of a node i is equal to the sum of its row in an adjacency matrix. Create a $N \times 1$ array (recall the dimensions of a matrix: *row* \times *column*) which stores the out-degree of each node in the input graph. Return the array of out-degrees.
- (3) Create a function `transitionProbabilities(adjacencyMatrix, outDegrees)` to generate and return the `transitionMatrix`. This is an $N \times N$ matrix with entries $p_{i,j}$ - which is the probability of the crawler moving: $i \rightarrow j$ defined as:

$$p_{i,j} = 0.90 \times A_{i,j}/O_i + 0.10/N$$

where A is the $(N \times N)$ adjacency matrix and O the (N) array of out-degrees. Compute and return the p matrix. Now you can proceed with computing the PageRank of each page in the graph:

I. PageRank by Simulation:

Using the functions you have implemented above, write a program `randomCrawler.py`. The program should take an integer `numTrials`, input by the user. The simulation starts the crawler at web-page 0. The crawler makes random transitions to other pages, keeping count of the number of times it visits each page.

Suppose the crawler is on page k . The transition probabilities from page k to every other page in the graph is given by the elements of the k^{th} row of the p (transition matrix).

To compute the page the crawler lands on:

- Pick a random number: $0 \leq r < 1$.
- Sum the probabilities till $\sum p[k][j] \geq r$.
- The crawler makes a transition to page j .
- Increment the count of visits to page j .

Repeat the above steps `numTrials` times. The elements of the count of visits array divided by `numTrials`, gives an approximation to the PageRank for each page in the graph.

II. PageRank using the Power Method:

For this part, do the familiar Markov Chain calculation:

- Start with an initial vector $x_0 = [1, 0, 0 \dots 0]^T$ (the superscript T indicates transpose). x_0 is a column vector ($N \times 1$ array).
- Carry out the computations: $x_{k+1} = p^T x_k$ for $k = 0, 1 \dots \text{numTrials}$.
- The final array $x_{\text{numTrials}}$ is the PageRank.

Problems:

- (1) Test your programs using `tiny.txt` with `numTrials = 100`. The output is provided alongside this document.
- (2) Repeat for `numTrials = 100000`. Make properly labelled histograms of the PageRanks for `tiny.txt` & `medium.txt`, for each of the above two methods.
- (3) Modify the graphs to ignore the effects of multiple links. That is, if there are multiple links from a page to another, count them as a single link. Do this with `tiny.txt` and `medium.txt`. Use the new transition matrix, compute the modified PageRank vector using the Power Method only. Include the histograms and comment on the changes you notice.
- (4) Finally, modify the file: `medium.txt` to add links to page 23 from every other page. Give the modified file a different name and save it. Run the Markov Chain calculation and make a properly labelled histogram of the PageRanks for the graph.

Note: Ideally, you will use a Jupyter notebook to present this project. You have used Jupyter notebooks in earlier labs. A reference for getting started with Jupyter notebooks is found at:

<https://www.dataquest.io/blog/jupyter-notebook-tutorial/>